

Yikai Wu

35 Olden St, Computer Science 431
Princeton, NJ 08540
USA
✉ yikai.wu@cs.princeton.edu

Education

Princeton University

Ph.D. student in Computer Science
Advisor: Prof. Sanjeev Arora

Princeton, NJ, USA

2022–present

Yale University

Ph.D. student in Computer Science

New Haven, CT, USA

2021–2022

Duke University

Bachelor of Science (Summa Cum Laude)
Double major in Computer Science and Mathematics, Minor in Physics
Overall GPA: 3.97/4.00
Member of honor society Phi Beta Kappa

Durham, NC, USA

2017–2021

Research Experience

LLM Agents for Algorithm Design and Optimization

Mentor: Dr. Ofir Press

Princeton University

Apr 2025–Jun 2025

- Contributed to the *AlgoTune* benchmark (154 expert-curated tasks) and baseline agent (*AlgoTuner*) that iteratively edits, compiles, profiles, and verifies code, achieving an average $1.72\times$ speedup over reference implementations (e.g., SciPy, scikit-learn, CVXPY).
- Create combinatorial-optimization tasks, provide sample solution codes with timing analysis, and evaluated LM-agent performance on these tasks across models.
- Paper [1] will be presented at the Conference on Neural Information Processing Systems (NeurIPS) 2025.

Machine learning for Combinatorial Optimization

Mentor: Prof. Sanjeev Arora

Princeton University

Oct 2023–Jan 2025

- We investigated ML approaches to the Maximum Independent Set (MIS), benchmarking recent AI-based methods against strong classical heuristics solvers. Across in-distribution random graphs and realistic settings, AI methods consistently lagged in solution quality and efficiency.
- Our serialization analysis shows non-backtracking AI effectively mimics simple degree-based greedy reasoning, explaining the gap to state-of-the-art classical methods.
- We documented these findings and recommendations in a position paper [2].

Differentially Private Optimization on Unitary Orbits

Mentor: Prof. Nisheeth Vishnoi

Yale University

Jun 2021–Feb 2022

- Consider the following optimization problem: Given $n \times n$ matrices A and Λ , maximize $\langle A, U\Lambda U \rangle$ where U varies over the unitary group $U(\Lambda)$. The optimization can be used in multiple matrix approximation problems such as PCA and rank- k approximation. We designed an efficient differentially private algorithm for this optimization where A is constructed using users' private data. We also provided upper and lower bounds for the utility of the optimization.
- Paper [3] was presented at the Conference on Learning Theory (COLT) 2022.

Hessian in Neural Networks

Duke University

Mentor: Prof. Rong Ge

May 2020–Aug 2021

- The Hessian of the loss function captures important properties of neural networks. We observed that the Hessian has several interesting structures which appear commonly in different neural networks. We explained these structures using Kronecker factorization. Our new understanding of the Hessian can be used with PAC-Bayes techniques to get better generalization bounds.
- Paper [4] on arXiv.

Differential Privacy for Multiple Analysts

Duke University

Mentor: Prof. Ashwin Machanavajjhala

Jan 2020–May 2021

- Differential privacy is the gold standard of privacy protection. Existing differentially private algorithms are designed for a single analyst and are problematic for multiple analysts. We formulated three criteria to decide whether an algorithm is good for multiple analysts. We demonstrated how existing algorithms fail to satisfy them. We also designed new differentially private algorithms which satisfy these criteria provably.
- Paper [5] was presented at the International Conference on Very Large Data Bases (VLDB) 2021. The preliminary version was presented at the Theory and Practice of Differential Privacy Workshop (TPDP) 2020.

Differential Privacy for Summation Queries

Duke University

Mentor: Prof. Ashwin Machanavajjhala

May 2018–Nov 2019

- Answering summation queries under differential privacy is a little understood, non-trivial problem. Traditional differentially private algorithms for these queries are data-independent and often result in large errors for some types of data. We proposed a data-dependent algorithm using truncation to effectively reduce the errors in the results, while providing the same level of privacy protection.
- Paper [6] was presented at the Theory and Practice of Differential Privacy Workshop (TPDP) 2019.

Quantum Information and Algorithms

Duke University

Mentor: Prof. Iman Marvian

May 2019–May 2020

- Quantum information applies modern physics knowledge to develop a new type of computers and communication devices. We investigated a type of quantum operators useful for ion-trap quantum computers. We also designed and analyzed algorithms to purify quantum states, which is useful for quantum communication and quantum error correction.
- In addition, I studied several research papers on quantum algorithms, quantum learning theory, and quantum complexity theory.

Publications

- [1] Ori Press, Brandon Amos, Haoyu Zhao, Yikai Wu, Samuel K Ainsworth, Dominik Krupke, Patrick Kidger, Touqir Sajed, Bartolomeo Stellato, Jisun Park, et al. Algotune: Can language models speed up general-purpose numerical programs? *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.
- [2] Yikai Wu, Haoyu Zhao, and Sanjeev Arora. Time to rethink ai for combinatorial optimization: Classical algorithms remain tough to match. 2025.
- [3] Oren Mangoubi, Yikai Wu, Satyen Kale, Abhradeep Thakurta, and Nisheeth K Vishnoi. Private matrix approximation and geometry of unitary orbits. In *Conference on Learning Theory*, pages 3547–3588. PMLR, 2022.
Presented at the Conference on Learning Theory (COLT) 2022.
- [4] Yikai Wu, Xingyu Zhu, Chenwei Wu, Annie Wang, and Rong Ge. Dissecting Hessian: Understanding common structure of Hessian in neural networks. *arXiv:2010.04261 [cs.LG]*, 2020.
Under review at the Conference on Neural Information Processing Systems (NeurIPS) 2022.
- [5] David Pujol, Yikai Wu, Brandon Fain, and Ashwin Machanavajjhala. Budget sharing for multi-analyst differential privacy. *Proceedings of the VLDB Endowment (PVLDB)*, 14(10): 1805–1817, 2021. doi: 10.14778/3467861.3467870.
Presented at the International Conference on Very Large Data Bases (VLDB) 2021.
- [6] Yikai Wu, David Pujol, Ios Kotsogiannis, and Ashwin Machanavajjhala. Answering summation queries for numerical attributes under differential privacy. *arXiv:1908.10268 [cs.DB]*, 2019.
Presented at the Theory and Practice of Differential Privacy Workshop (TPDP) 2019.

Teaching Assistant

COMPSCI 590.07: Computational Microeconomics (Graduate) <i>Instructor: Prof. Vincent Conitzer</i>	Duke University <i>Fall 2020</i>
COMPSCI 230: Discrete Mathematics for Computer Science <i>Instructor: Prof. Kamesh Munagala</i>	Duke University <i>Spring 2018</i>

Honors and Awards

Computing Research Association (CRA) Outstanding Undergraduate Researcher Honorable Mention	2020
Duke University Faculty Scholar Nomination (Top 2 in Computer Science department)	2020
Mathematical Contest in Modeling (MCM) Meritorious Winner	2019
Duke University Mathematics Student Award The Karl Menger Award	2018

The International Collegiate Programming Contest (ICPC)

Mid-Atlantic Regional Ranked 4th

2017, 2018

William Lowell Putnam Mathematical Competition

Ranked 142.5 (Top 4 at Duke Univeristy)

2017

Computer skills

Python, PyTorch, Tensorflow, C, C++, R, Mathematica